

1.

query: select \* from section;

chosen plan: "Seq Scan on section (cost=0.00..2.00 rows=100 width=28) (actual time=0.014..0.027 rows=100 loops=1)"

"Planning time: 0.063 ms"

"Execution time: 0.058 ms"

reason: simple select statement uses seq scanning in file.

2.

query: create index nam\_idx on student(name);

explain analyze select \* from student where name='Schrefl';

chosen plan: "Bitmap Heap Scan on student (cost=4.31..14.16 rows=4 width=24) (actual time=0.037..0.040 rows=4 loops=1)"

" Recheck Cond: ((name)::text = 'Schrefl'::text)"

" Heap Blocks: exact=3"

" -> Bitmap Index Scan on nam\_idx (cost=0.00..4.31 rows=4 width=0) (actual time=0.028..0.029 rows=4 loops=1)"

" Index Cond: ((name)::text = 'Schrefl'::text)"

"Planning time: 1.328 ms"

"Execution time: 0.079 ms"

reason: searching for some value after creating an index does an Bitmap heap scan in the created index.

3. query: select \* from student where name='rumat' and id='0000';

chosen plan: "Index Scan using name\_idx on student (cost=0.28..8.30 rows=1 width=24) (actual time=0.022..0.022 rows=0 loops=1)"

" Index Cond: ((name)::text = 'rumat'::text)"

" Filter: ((id)::text = '0000'::text)"

"Planning time: 0.779 ms"

"Execution time: 0.045 ms"

reason: searching for two values simultaneously does an index scan and then filters for the other attribute.

4.

query: select \* from takes, student

where (takes.id, takes.course\_id) = ('83314', '960') and student.id = '0000';

chosen plan: "Nested Loop (cost=0.56..16.61 rows=1 width=48) (actual time=0.069..0.069 rows=0 loops=1)"

" -> Index Scan using takes\_pkey on takes (cost=0.29..8.31 rows=1 width=24) (actual time=0.043..0.047 rows=2 loops=1)"

" Index Cond: (((id)::text = '83314'::text) AND ((course\_id)::text = '960'::text))"

```
" -> Index Scan using student_pkey on student (cost=0.28..8.29 rows=1 width=24) (actual
time=0.009..0.009 rows=0 loops=2)"
"    Index Cond: ((id)::text = '0000'::text)"
"Planning time: 1.030 ms"
"Execution time: 0.131 ms"
```

reason: searching for a tuple in a relation does multiple index scans inside a nested loop.

5.

query:

```
select * from (select * from takes order by id) as temp_s , (select * from student order by id)
as temp_t
where temp_t.id = temp_s.id;
```

chosen plan:

```
"Merge Join (cost=0.56..3011.55 rows=30000 width=48) (actual time=0.031..46.859
rows=30000 loops=1)"
"  Merge Cond: ((student.id)::text = (takes.id)::text)"
"    -> Index Scan using student_pkey on student (cost=0.28..130.27 rows=2000 width=24)
(actual time=0.013..1.680 rows=2000 loops=1)"
"    -> Materialize (cost=0.29..2481.28 rows=30000 width=24) (actual time=0.011..30.048
rows=30000 loops=1)"
"      -> Index Scan using takes_pkey on takes (cost=0.29..2106.28 rows=30000 width=24)
(actual time=0.009..24.540 rows=30000 loops=1)"
"Planning time: 0.458 ms"
"Execution time: 48.936 ms"
```

reason: using order by clause with join query uses merge join (as given in the hint)

6.

after adding limit 10 to the previous query, the plan:

```
"Limit (cost=0.56..1.57 rows=10 width=48) (actual time=0.032..0.052 rows=10 loops=1)"
"  -> Merge Join (cost=0.56..3011.55 rows=30000 width=48) (actual time=0.032..0.049
rows=10 loops=1)"
"    Merge Cond: ((student.id)::text = (takes.id)::text)"
"      -> Index Scan using student_pkey on student (cost=0.28..130.27 rows=2000
width=24) (actual time=0.017..0.017 rows=1 loops=1)"
"      -> Materialize (cost=0.29..2481.28 rows=30000 width=24) (actual time=0.010..0.023
rows=10 loops=1)"
"        -> Index Scan using takes_pkey on takes (cost=0.29..2106.28 rows=30000
width=24) (actual time=0.007..0.018 rows=10 loops=1)"
"Planning time: 3.247 ms"
"Execution time: 0.100 ms"
```

since the algorithm did not change,

a different query where it changes:

select \* from prereq natural join takes order by course\_id;

chosen plan before adding limit 10:

```
"Merge Join (cost=2759.39..3191.53 rows=28481 width=28) (actual time=20.703..26.118
rows=16860 loops=1)"
"  Merge Cond: ((prereq.course_id)::text = (takes.course_id)::text)"
"    -> Sort (cost=5.32..5.57 rows=100 width=8) (actual time=0.066..0.075 rows=100
loops=1)"
"      Sort Key: prereq.course_id"
"      Sort Method: quicksort Memory: 29kB"
"        -> Seq Scan on prereq (cost=0.00..2.00 rows=100 width=8) (actual time=0.004..0.009
rows=100 loops=1)"
"      -> Sort (cost=2750.90..2825.90 rows=30000 width=24) (actual time=20.496..22.001
rows=34008 loops=1)"
"        Sort Key: takes.course_id"
"        Sort Method: quicksort Memory: 3112kB"
"          -> Seq Scan on takes (cost=0.00..520.00 rows=30000 width=24) (actual
time=0.004..4.811 rows=30000 loops=1)"
"Planning time: 4.190 ms"
"Execution time: 26.789 ms"
```

chosen plan after adding limit 10:

```
"Limit (cost=0.43..9.17 rows=10 width=28) (actual time=3.165..3.199 rows=10 loops=1)"
"  -> Nested Loop (cost=0.43..24904.39 rows=28481 width=28) (actual time=3.164..3.197
rows=10 loops=1)"
"    -> Index Only Scan using prereq_pkey on prereq (cost=0.14..13.64 rows=100
width=8) (actual time=1.953..1.954 rows=2 loops=1)"
"      Heap Fetches: 2"
"    -> Index Scan using takes_pkey on takes (cost=0.29..245.38 rows=353 width=24)
(actual time=0.603..0.619 rows=5 loops=2)"
"      Index Cond: ((course_id)::text = (prereq.course_id)::text)"
"Planning time: 2.702 ms"
"Execution time: 3.222 ms"
```

hence after adding limit 10, the algorithm changes from merge join to index scan inside nested loop.

Reason: takes is a small relation, hence the nested is preferred for lesser time than merge join.

7.

time taken for creating index is 142 ms

time taken for dropping index is 2 ms

8.

8.1

begin;

Query returned successfully with no result in 15 msec

8.2

"Delete on course (cost=0.00..4.50 rows=1 width=6) (actual time=0.061..0.061 rows=0 loops=1)"

" -> Seq Scan on course (cost=0.00..4.50 rows=1 width=6) (actual time=0.012..0.035 rows=1 loops=1)"

" Filter: ((course\_id)::text = '400'::text)"

" Rows Removed by Filter: 199"

"Planning time: 1.406 ms"

"Trigger for constraint section\_course\_id\_fkey on course: time=1.397 calls=1"

"Trigger for constraint prereq\_course\_id\_fkey on course: time=1.587 calls=1"

"Trigger for constraint prereq\_prereq\_id\_fkey on course: time=0.380 calls=1"

"Trigger for constraint teaches\_course\_id\_fkey on section: time=2.255 calls=2"

"Trigger for constraint takes\_course\_id\_fkey on section: time=9.364 calls=2"

"Execution time: 15.075 ms"

8.3

rollback;

Query returned successfully with no result in 11 msec.

8.4

create index ind1 on section(course\_id);

Query returned successfully with no result in 22 msec.

create index ind2 on prereq(course\_id);

Query returned successfully with no result in 23 msec.

create index ind3 on prereq(prereq\_id);

Query returned successfully with no result in 23 msec.

create index ind4 on teaches(course\_id, sec\_id, semester, year);

Query returned successfully with no result in 22 msec.

create index ind5 on takes(course\_id, sec\_id, semester, year);

Query returned successfully with no result in 142 msec.

8.5

delete index ind1;

Query returned successfully with no result in 2 msec.

Delete index ind2;

Query returned successfully with no result in 0 msec.

delete index ind3;

Query returned successfully with no result in 1 msec.

delete index ind4;

Query returned successfully with no result in 12 msec.

delete index ind5;

Query returned successfully with no result in 8 msec.

8.6

while executing drop, the time is lesser because we postgresql does not read the whole data, it can only drop the allocated data directly.